

Data Science and Machine Learning in Python

Stephan Weyers

Part 1: Data Science

	Date	Topics covered
1	Apr 13 th	Course introduction Data Science motivation How to use Jupyter Notebook Python types and lists Loops, if/else, functions
2	Apr 20 th	Python tuples, lists, dictionaries Functions Numpy basics, operations Image processing
3	Apr 27 th	Pandas Series, DataFrame Pandas basic operations Import/export files
4	May 4 th	Principles of data visualization Data cleaning and preparation Join, combine and reshape data
5	May 11 th	Volkswahl Bund dataset Data visualization in Python How to write Data Science reports Data aggregation and grouping

Part 2: Machine Learning

	Date	Topics covered
6	Jun 1 st	Introduction to supervised learning Classification and regression scikit-learn k-Nearest Neighbors Linear regression (ridge and lasso)
7	Jun 8 th	Linear classification models Decision trees Random forests and gradient boosting
8	Jun 15 th	Kernel support vector machines Neural networks
9	Jun 22 nd	Introduction to unsupervised learning Preprocessing and scaling Dimensionality reduction Principal component analysis
10	Jun 29 th	k-means clustering Hierarchical clustering DBSCAN
11	Jul 6 th	Representing data Engineering features Model evaluation and improvement Text data analysis

Deadlines for Submission and Distribution of Grading

Student task	Deliverables	Deadline	Work	Share of grade
W01 Assignment	Code and results	Apr 26 th	Team A	5.0%
W02 Case Study	Code / presentation slides	May 22 nd	Team B	18.0%
W02 Case Study	Peer review*	May 31 st	Individual	2.0%
W03 Assignment	Code and results	May 29 th	Team B	5.0%
W04 Assignment	Code and results	Jun 12 th	Team C	10.0%
W05 Assignment	Code and results	Jun 28 th	Team D	7.0%
W06 Assignment	Code and results	Jul 8 th	Team D	13.0%
W07 Case Study	Code / presentation slides	Jul 17 th	Team D	22.0%
W07 Case Study	Peer review*	Jul 31 st	Individual	3.0%
DataCamp 1	Finish course	May 9 th	Individual	2.5%
DataCamp 2	Finish course	May 30 th	Individual	2.5%
DataCamp 3	Finish course	Jun 20 th	Individual	2.5%
DataCamp 4	Finish course	Jul 11 th	Individual	2.5%

* Peer review is mandatory. Quality of peer review itself is graded. Not providing peer review at all would result in high point deduction

Agenda for online lecture 11

Session	Topic	Mode	Materials used	Minutes	End
14:30-16:00	Organizational questions	Q&A		10	14:40
	OCEAN Big Five	Lecture / Q&A	Lecture 10d notebook	10	14:50
	Text clusters – introduction	Lecture / Q&A	Lecture 11a notebook	5	14:55
	Find and label text clusters	Team work in break-out rooms	Lecture 11a notebook	25	15:20
	Text cluster results	Discussion in main room	Lecture 11a notebook	10	15:30
	Model tuning / evaluation	Lecture / Q&A	Lecture slides	15	15:45
	Course evaluation	Individual work	Evaluation form	10	15:55
16:10-17:40	Example churn – intro	Lecture / Q&A	Lecture 11b notebook	20	16:30
	Example churn – Exercise	Joint discussion in main room	Lecture 11b notebook	40	17:10
	Course evaluation	Team work in break-out rooms	2 stars + 2 wishes	25	17:30
	Farewell			5	17:40
17:50-19:20	Optional	Optional	Optional		

Dataset Source: kaggle (<https://www.kaggle.com/tunguz/big-five-personality-test>)

Attribute Information:

The scale was labeled 1=Disagree, 3=Neutral, 5=Agree

E - Surgency or Extraversion

EXT1 I am the life of the party. EXT2 I don't talk a lot. EXT3 I feel comfortable around people. EXT4 I keep in the background. EXT5 I start conversations. EXT6 I have little to say. EXT7 I talk to a lot of different people at parties. EXT8 I don't like to draw attention to myself. EXT9 I don't mind being the center of attention. EXT10 I am quiet around strangers.

N - Emotional Stability or (not) Neuroticism

EST1 I get stressed out easily. EST2 I am relaxed most of the time. EST3 I worry about things. EST4 I seldom feel blue. EST5 I am easily disturbed. EST6 I get upset easily. EST7 I change my mood a lot. EST8 I have frequent mood swings. EST9 I get irritated easily. EST10 I often feel blue.

A - Agreeableness

AGR1 I feel little concern for others. AGR2 I am interested in people. AGR3 I insult people. AGR4 I sympathize with others' feelings. AGR5 I am not interested in other people's problems. AGR6 I have a soft heart. AGR7 I am not really interested in others. AGR8 I take time out for others. AGR9 I feel others' emotions. AGR10 I make people feel at ease.

C - Conscientiousness

CSN1 I am always prepared. CSN2 I leave my belongings around. CSN3 I pay attention to details. CSN4 I make a mess of things. CSN5 I get chores done right away. CSN6 I often forget to put things back in their proper place. CSN7 I like order. CSN8 I shirk my duties. CSN9 I follow a schedule. CSN10 I am exacting in my work.

O - Openness to experience or Intellect or Imagination

OPN1 I have a rich vocabulary. OPN2 I have difficulty understanding abstract ideas. OPN3 I have a vivid imagination. OPN4 I am not interested in abstract ideas. OPN5 I have excellent ideas. OPN6 I do not have a good imagination. OPN7 I am quick to understand things. OPN8 I use difficult words. OPN9 I spend time reflecting on things. OPN10 I am full of ideas.

df

874434 rows x 50 columns

6

```
from sklearn.cluster import KMeans  
kmeans = KMeans(n_clusters=5, random_state=10)  
kmeans.fit(df)
```

```
KMeans(n_clusters=5, random_state=10)
```

```
print(df.loc[:, "EXT1"].groupby(kmeans.labels_).count())  
p_mean = df.groupby(kmeans.labels_).mean().transpose()  
import seaborn as sns  
plt.figure(figsize = (10,20))  
ax = sns.heatmap(p_mean, yticklabels=df.columns.values, annot=True, fmt=".2f", cmap="PiYG")  
# Observation: The five types of variables really mostly appear to be "packages"
```


```
0    178189  
1    222885  
2    181026  
3    166348  
4    125986  
Name: EXT1, dtype: int64
```

OCEAN Big Five – kmeans

EXT1	0.76	0.37	-0.59	-0.29	-0.49
EXT2	0.79	0.37	-0.52	-0.34	-0.56
EXT3	0.57	0.68	-0.69	-0.31	-0.60
EXT4	0.72	0.47	-0.70	-0.29	-0.46
EXT5	0.72	0.55	-0.57	-0.41	-0.63
EXT6	0.61	0.42	-0.31	-0.62	-0.33
EXT7	0.76	0.49	-0.63	-0.35	-0.57
EXT8	0.68	0.22	-0.53	-0.22	-0.29
EXT9	0.71	0.29	-0.55	-0.37	-0.25
EXT10	0.65	0.50	-0.67	-0.32	-0.42
EST1	-0.22	0.63	-0.67	-0.24	0.48
EST2	-0.09	0.48	-0.56	-0.14	0.29
EST3	-0.21	0.49	-0.63	-0.11	0.48
EST4	-0.15	0.45	-0.52	0.01	0.14
EST5	-0.22	0.55	-0.43	-0.27	0.31
EST6	-0.30	0.68	-0.62	-0.25	0.45
EST7	-0.43	0.73	-0.57	-0.17	0.37
EST8	-0.40	0.75	-0.61	-0.18	0.36
EST9	-0.28	0.73	-0.49	-0.22	0.10
EST10	-0.13	0.75	-0.78	-0.11	0.13
AGR1	0.13	0.31	0.25	-0.28	-0.73
AGR2	0.48	0.47	-0.01	-0.39	-0.97
AGR3	-0.29	0.45	0.06	-0.02	-0.45
AGR4	0.24	0.34	0.37	-0.19	-1.23
AGR5	0.29	0.37	0.20	-0.23	-1.05
AGR6	0.17	0.15	0.37	-0.01	-1.03
AGR7	0.42	0.51	0.02	-0.34	-1.06
AGR8	0.22	0.39	0.14	-0.23	-0.89
AGR9	0.30	0.30	0.35	-0.23	-1.15
AGR10	0.32	0.52	-0.17	-0.37	-0.63

AGR1	0.13	0.31	0.25	-0.28	-0.73
AGR2	0.48	0.47	-0.01	-0.39	-0.97
AGR3	-0.29	0.45	0.06	-0.02	-0.45
AGR4	0.24	0.34	0.37	-0.19	-1.23
AGR5	0.29	0.37	0.20	-0.23	-1.05
AGR6	0.17	0.15	0.37	-0.01	-1.03
AGR7	0.42	0.51	0.02	-0.34	-1.06
AGR8	0.22	0.39	0.14	-0.23	-0.89
AGR9	0.30	0.30	0.35	-0.23	-1.15
AGR10	0.32	0.52	-0.17	-0.37	-0.63
CSN1	-0.36	0.50	-0.11	-0.22	0.07
CSN2	-0.51	0.42	-0.15	0.10	0.06
CSN3	-0.20	0.32	0.17	-0.43	0.04
CSN4	-0.45	0.70	-0.38	-0.14	0.13
CSN5	-0.37	0.54	-0.21	-0.00	-0.13
CSN6	-0.48	0.54	-0.19	-0.04	0.06
CSN7	-0.33	0.29	0.08	-0.13	-0.00
CSN8	-0.29	0.57	-0.18	-0.21	-0.06
CSN9	-0.28	0.43	-0.06	-0.05	-0.20
CSN10	-0.17	0.36	0.04	-0.41	0.08
OPN1	0.14	0.14	0.22	-0.85	0.34
OPN2	0.09	0.26	0.14	-0.87	0.37
OPN3	0.30	0.00	0.39	-0.80	0.07
OPN4	0.14	0.16	0.24	-0.76	0.19
OPN5	0.27	0.31	0.01	-0.91	0.28
OPN6	0.24	0.18	0.24	-0.83	0.10
OPN7	0.05	0.32	0.02	-0.77	0.34
OPN8	0.19	-0.04	0.25	-0.70	0.38
OPN9	0.05	0.01	0.43	-0.55	0.02
OPN10	0.34	0.23	0.20	-1.05	0.20
	0	1	2	3	4



```
df2 = pd.DataFrame()
df2["EXT"] = df.loc[:, "EXT1": "EXT10"].sum(axis=1)
df2["EST"] = df.loc[:, "EST1": "EST10"].sum(axis=1)
df2["AGR"] = df.loc[:, "AGR1": "AGR10"].sum(axis=1)
df2["CSN"] = df.loc[:, "CSN1": "CSN10"].sum(axis=1)
df2["OPN"] = df.loc[:, "OPN1": "OPN10"].sum(axis=1)
df2["Cluster"] = kmeans.labels_
p_mean = df2.groupby("Cluster").mean().transpose()
import seaborn as sns
plt.figure(figsize = (10,5))
ax = sns.heatmap(p_mean, yticklabels=df2.columns.values, annot=True, fmt=".2f", cmap="PiYG")
```





```
from sklearn.decomposition import PCA
pca = PCA(n_components=2)
pca.fit(df)
projected = pca.transform(df)

|
colors = ["blue", "orange", "green", "red", "purple"]
plt.figure(figsize=(10,10))
for i in range(5):
    Z = projected[kmeans.labels_ == i]
    plt.scatter(Z[:,0], Z[:,1], color=colors[i], s=9, label=i)
plt.legend(loc="upper right")
plt.show()
```



```
from sklearn.cluster import AgglomerativeClustering  
agg = AgglomerativeClustering(n_clusters=5)  
agg.fit(df)
```

MemoryError

Traceback (most recent call last)

MemoryError: Unable to allocate 2.78 TiB for an array with shape (382316972961,) and data type float64

sentences[12]

'crude oil prices back above \$50 cold weather across parts of the united states and much of europe has pushed us crude oil prices above \$50 a barrel for the first time in almost three months. freezing temperatures and heavy snowfall have increased demand for heating fuel in the us where stocks are low. fresh falls in the value of the dollar helped carry prices above the \$50 mark for the first time since november. a barrel of us crude oil closed up \$2.80 to \$51.15 in new york on tuesday. opec members said on tuesday that it saw no reason to cut its output. although below last year's peak of \$55.67 a barrel which was reached in october prices are now well above 2004's average of \$41.48. brent crude also rose in london trading adding \$1.89 to \$48.62 at the close. much of western europe and the north east of america has been shivering under unseasonably low temperatures in recent days. the decline in the us dollar to a five-week low against the euro has also served to inflate prices. the dollar moved sharply overnight and oil is following it said chris furness senior market strategist at 4cast. if the dollar continues to weaken oil will be obviously higher. several opec members said a cut in production was unlikely citing rising prices and strong demand for oil from asia. i agree that we do not need to cut supply if the prices are as much as this fathi bin shatwan libya's oil minister told reuters. i do not think we need to cut unless the prices are falling below \$35 a barrel he added. opec closely watches global stocks to ensure that there is not an excessive supply in the market. the arrival of spring in the northern hemisphere will focus attention on stockpiles of us crude and gasoline which are up to 9% higher than at this time last year. heavy stockpiles could help force prices lower when demand eases.\n'

sentences[1215]

'wru proposes season overhaul the welsh rugby union wants to restructure the northern hemisphere season into four separate blocks. the season would start with the celtic league in october followed by the heineken cup in february and march and the six nations moved to april and may. after a nine week break the wrw then proposes a two-month period of away and home international matches. wrw chairman david pickering said the structure would end problems of player availability for club and country. he added: we feel sure that spectator interest would respond to the impetus of high intensity rugby being played continuously rather than the fragmented timetable currently in operation. equally we suspect that the sponsors would prefer the sustained interest in a continuous tournament and hopefully the broadcasters would also enjoy increased exposure. moving the six nations from its traditional february beginning should also ensure better weather conditions and stimulate greater interest in the games and generally provide increased skills and competition and attract greater spectator viewing pickering argued. the plan will be put before the international rugby board next month where four other plans drawn up by independent consultants for a global integrated season will also be discussed. pickering added: it's very early days and there are a number of caveats associated with it - not least the revenue from the broadcasters which is extremely important. we've got a good plan and one which should be judged on its merits.\n'

```
## full_remove takes a string x and a list of characters removal_list
## returns x with all the characters in removal_list replaced by ' '
def full_remove(x, removal_list):
    for w in removal_list:
        x = x.replace(w, ' ')
    return x
## Remove digits
digits = [str(x) for x in range(10)]
digit_less = [full_remove(x, digits) for x in sentences]
## Remove punctuation
punc_less = [full_remove(x, list(string.punctuation)) for x in digit_less]
## Make everything lower-case
sents_lower = [x.lower() for x in punc_less]
## Define our stop words
stop_set = set(['the', 'a', 'an', 'i', 'he', 'she', 'they', 'to', 'of', 'it', 'from',
    "and", "in", "is", "for", "that", "on", "was", "be", "with",
    "as", "has", "have", "at", "are", "but", "will", "by",
    "this", "which", 'then', 'him', 'going', 'any', 'while', 'before',
    'because', 'should', 'many', 'three', 'very', 'made', 'such', 'get',
    'told', 'being', 'just', 'best', 'no', 'into', 'some', 'what',
    'so', 'now', 'two', 'when', 'over', 'could', 'if', 'you', 'all', 'than',
    'or', 'can', 'about', 'there', 'one', 'us', 'new', 'who', 'also', 'were',
    'its', 'their', 'been', 'had', 'would', 'his', 'not', 'we', 'said', 'after',
    'out', 'more', 'up', 'first', 'last', 'like', 'make', 'only', 'do', 'other', 'her',
    'year', 'years', 'them', 'against', 'back', 'next', 'bbc', 'well', 'set', 'number',
    'take', 'most', 'way', 'added', 'may', 'says', 'my', 'our', 'off', 'good',
    'how', 'down', 'still', 'those', 'much', 'uk', 'england', 'go', 'since', 'say'])
## Remove stop words
sents_split = [x.split() for x in sents_lower]
sents_processed = [" ".join(list(filter(lambda a: a not in stop_set, x))) for x in sents_split]
```



```
n = 12
print(sentences[n])
print(sents_processed[n])
```

crude oil prices back above \$50 cold weather across parts of the united states and much of europe has pushed us crude oil prices above \$50 a barrel for the first time in almost three months. freezing temperatures and heavy snowfall have increased demand for heating fuel in the us where stocks are low. fresh falls in the value of the dollar helped carry prices above the \$50 mark for the first time since november. a barrel of us crude oil closed up \$2.80 to \$51.15 in new york on tuesday. opec members said on tuesday that it saw no reason to cut its output. although below last year's peak of \$55.67 a barrel which was reached in october prices are now well above 2004's average of \$41.48. brent crude also rose in london trading adding \$1.89 to \$48.62 at the close. much of western europe and the north east of america has been shivering under unseasonably low temperatures in recent days. the decline in the us dollar to a five-week low against the euro has also served to inflate prices. the dollar moved sharply overnight and oil is following it said chris furness senior market strategist at 4cast. if the dollar continues to weaken oil will be obviously higher. several opec members said a cut in production was unlikely citing rising prices and strong demand for oil from asia. i agree that we do not need to cut supply if the prices are as much as this fathi bin shatwan libya's oil minister told reuters. i do not think we need to cut unless the prices are falling below \$35 a barrel he added. opec closely watches global stocks to ensure that there is not an excessive supply in the market. the arrival of spring in the northern hemisphere will focus attention on stockpiles of us crude and gasoline which are up to 9% higher than at this time last year. heavy stockpiles could help force prices lower when demand eases.

crude oil prices above cold weather across parts united states europe pushed crude oil prices above barrel time almost month's freezing temperatures heavy snowfall increased demand heating fuel where stocks low fresh falls value dollar helped carry prices above mark time november barrel crude oil closed york tuesday opec members tuesday saw reason cut output although below's peak barrel reached october prices above's average brent crude rose london trading adding close western europe north east america shivering under unseasonably low temperatures recent days decline dollar five week low euro served inflate prices dollar moved sharply overnight oil following chris furness senior market strategist cast dollar continues weaken oil obviously higher several opec members cut production unlikely citing rising prices strong demand oil asia agree need cut supply prices fathi bin shatwan libya's oil minister reuters think need cut unless prices falling below barrel opec closely watches global stocks ensure excessive supply market arrival spring northern hemisphere focus attention stockpiles crude gasoline higher time heavy stockpiles help force prices lower demand eases

```
## Transform to bag of words representation.  
from sklearn.feature_extraction.text import CountVectorizer  
vectorizer = CountVectorizer(min_df = 20, max_df=100000, max_features = None)  
data_features = vectorizer.fit_transform(sents_processed)  
X = data_features.toarray()  
vocabulary = vectorizer.get_feature_names()  
df = pd.DataFrame(X, columns=vocabulary.get_feature_names())  
X.shape, df.shape
```

```
((2225, 3258), (2225, 3258))
```

```
# Show top 100 words that appear most often  
df.sum(axis=0).sort_values(ascending=False).head(50).index
```

```
Index(['mr', 'people', 'time', 'world', 'government', 'bn', 'film', 'game',  
      'music', 'labour', 'market', 'company', 'home', 'election', 'party',  
      'games', 'win', 'work', 'firm', 'second', 'top', 'blair', 'show', 'won',  
      'think', 'week', 'use', 'million', 'part', 'play', 'technology',  
      'minister', 'high', 'public', 'want', 'between', 'under', 'mobile',  
      'see', 'british', 'did', 'five', 'country', 'used', 'european', 'tv',  
      'players', 'through', 'news', 'end'],  
      dtype='object')
```


term frequency–inverse document frequency (tf–idf): Give high weight terms that appear often in a particular document, but not in many documents in the corpus

n-grams:

```
bards_words:
['The fool doth think he is wise,',
 'but the wise man knows himself to be a fool']

Vocabulary:
['be fool', 'but the', 'doth think', 'fool doth', 'he is', 'himself to',
 'is wise', 'knows himself', 'man knows', 'the fool', 'the wise',
 'think he', 'to be', 'wise man']
```

Lemmatization and Stemming:

```
compare_normalization(u"Our meeting today was worse than yesterday, "
                      "I'm scared of meeting the clients tomorrow.")
```

Lemmatization:

```
['our', 'meeting', 'today', 'be', 'bad', 'than', 'yesterday', ',', 'i', 'be',
 'scared', 'of', 'meet', 'the', 'client', 'tomorrow', '.']
```

Stemming:

```
['our', 'meet', 'today', 'wa', 'wors', 'than', 'yesterday', ',', 'i', 'm',
 'scare', 'of', 'meet', 'the', 'client', 'tomorrow', '.']
```

```
# Perform kmeans clustering
from sklearn.cluster import KMeans
kmeans = KMeans(n_clusters=4, random_state=20)
kmeans.fit(X)
# Show number of cases per cluster
df["CLUSTER"] = kmeans.labels_
df.iloc[:, -2:].groupby("CLUSTER").count().transpose()
```

	CLUSTER	0	1	2	3
--	---------	---	---	---	---

zealand	5	1942	277	1
---------	---	------	-----	---

```
# Display top 50 words appearing most often in cluster n
# as well as the m-th sentence in cluster n
n = 2
m = 1
print(df[df["CLUSTER"]==n].drop("CLUSTER", axis=1).sum(axis=0).sort_values(ascending=False).head(50).index)
i = df[df["CLUSTER"]==n].index[m]
sentences[i]
```

```
Index(['mr', 'labour', 'election', 'blair', 'party', 'government', 'people',
      'brown', 'minister', 'howard', 'prime', 'tory', 'tax', 'public',
      'plans', 'leader', 'chancellor', 'tories', 'britain', 'campaign',
      'secretary', 'home', 'general', 'bn', 'tony', 'time', 'lib', 'michael',
      'country', 'under', 'between', 'world', 'budget', 'spokesman',
      'kennedy', 'liberal', 'saying', 'lord', 'part', 'mps', 'eu', 'think',
      'vote', 'issue', 'british', 'conservative', 'want', 'claims',
      'spending', 'did'],
      dtype='object')
```

'howard hits back at mongrel jibe michael howard has said a claim by peter hain that the tory leader is acting like an attack mongrel shows labour is rattled by the opposition. in an upbeat speech to his party's spring conference in brighton he said labour's campaigning tactics proved the tories were hitting home. mr hain made the claim about tory tactics in the anti-terror bill debate. something tells me that someone somewhere out there is just a little bit rattled mr howard said. mr hain leader of the commons told bbc radio four's today programme that mr howard's stance on the government's anti-terrorism legislation was putting the country at risk. he then accused the tory leader of behaving like an attack mongrel and playing opposition for opposition sake. mr howard told his party that labour would do anything say anything claim anything

```
# Perform hierarchical clustering
from sklearn.cluster import AgglomerativeClustering
agg = AgglomerativeClustering(n_clusters=4, linkage="complete")
agg.fit(X)
# Show number of cases per cluster
df["CLUSTER"] = agg.labels_
df.iloc[:, -2:].groupby("CLUSTER").count().transpose()
```



CLUSTER	0	1	2	3
zealand	2	1	2221	1

```
# Display top 50 words appearing most often in cluster n
# as well as the m-th sentence in cluster n
n = 1
m = 0
print(df[df["CLUSTER"]==n].drop("CLUSTER", axis=1).sum(axis=0).sort_values(ascending=False).head(50).index)
i = df[df["CLUSTER"]==n].index[m]
sentences[i]
```

```
Index(['people', 'time', 'mr', 'world', 'film', 'game', 'bn', 'music',
      'market', 'company', 'government', 'games', 'firm', 'second', 'top',
      'win', 'show', 'won', 'million', 'technology', 'mobile', 'play', 'use',
      'players', 'tv', 'high', 'home', 'work', 'week', 'sales', 'five', 'end',
      'half', 'used', 'part', 'however', 'through', 'net', 'four', 'want',
      'growth', 'six', 'european', 'both', 'group', 'think', 'see', 'old',
      'between', 'did'],
      dtype='object')
```

'tv future in the hands of viewers with home theatre systems plasma high-definition tvs and digital video recorders moving into the living room the way people watch tv will be radically different in five years time. that is according to an expert panel which gathered at the annual consumer electronics show in las vegas to discuss how these new technologies will impact one of our favourite pastimes. with the us leading the trend programmes and other content will be delivered to viewers via home networks through cable satellite telecoms companies and broadband service providers to front rooms and portable devices. one of the most talked-about technologies of ces has been digital and personal video recorders (dvr and pvr). these set-top boxes like the us s tivo and the uk s sky+ system allow people to record store play pause and forward wind tv programmes when they want. essentially the technology allows for much more personalised tv. they are also being built-in t


Text Cluster – Elbow graph




Find a good cluster solution and give appropriate names to the resulting clusters!

- Change the number of clusters in kmeans and hierarchical clustering
- Change the linkage method
- Potentially change the stopwords or the parameters in CountVectorizer
- Explore the clusters by looking at the top 50 words in each cluster and sample sentences

How many distinct meaningful clusters do you find?





Examples

Scaling / Rescaling

- Standardize (mean 0, variance 1)
- Min/Max scaling
- Logarithm, square root, inverse,...

Categorization

- Binning of metric variables
- Transform zip code to area (East, West, South, North)
- Combine two categorical variables (e.g. area with income group)

Ratios

- Sales above/below average
- Average sales per customer
- Purchases per view
- Share of sales in category X

Segmentation / dimensionality reduction

- Result of cluster analysis (k-means, hierarchical clustering,...) as new variable
- New variables through principal component analysis, factor analysis,...

Balanced data



Imbalanced data



Undersampling



Oversampling



Question for discussion

- Find examples for balanced and imbalanced datasets.
- What are advantages and disadvantages of undersampling and oversampling?

predict_proba

[0.78 , 0.22]

[0.99 , 0.01]

[0.41 , 0.59]

[0.02 , 0.98]

[0.67 , 0.33]

[0.11 , 0.89]

Confusion matrix

		Predicted churn	
		No	Yes
Actual churn	No	True Negative	False Positive
	Yes	False Negative	True Positive

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}$$

$$\text{Precision} = \frac{TP}{TP + FP}$$

$$\text{Recall} = \frac{TP}{TP + FN}$$

$$F_1 = \frac{2 \cdot \text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}}$$

3333 rows (customers)

21 variables:


- State: customer residence, e.g. OH or NJ
- Account Length: number of days this account has been active
- Area Code: 3-digit area code of customer's phone number
- Phone: remaining seven-digit phone number
- Int'l Plan: customer has international calling plan: yes/no
- VMail Plan: customer has voice mail feature: yes/no
- VMail Message: average number of voice mail messages per month
- Day Mins: total number of calling minutes used during the day
- Day Calls: total number of calls placed during the day
- Day Charge: billed cost of daytime calls
- Eve Mins, Eve Calls, Eve Charge: same for evening calls
- Night Mins, Night Calls, Night Charge: same for night calls
- Intl Mins, Intl Calls, Intl Charge: same for international calls
- CustServ Calls: number of calls placed to Customer Service
- Churn: customer left the service: true/false

```
print(telco['Churn'].value_counts())
```


```
no      2850  
yes      483  
Name: Churn, dtype: int64
```

Example Churn – Distributions

```
sns.displot(data=telco, x='Account_Length', kind="hist", height=2, aspect=4)  
plt.show()
```




```
sns.displot(data=telco, x='Intl_Mins', kind="hist", height=2, aspect=4)  
plt.show()
```




Example Churn – Boxplots

```
plt.figure(figsize=(6, 10))  
sns.boxplot(x = 'Churn', y = 'Day_Mins', data = telco)  
plt.show()
```



```
plt.figure(figsize=(6, 10))  
sns.boxplot(x = 'Churn', y = 'CustServ_Calls', data = telco)  
plt.show()
```



```
telco.dtypes
```

```
Account_Length      int64
Vmail_Message       int64
Day_Mins            float64
Eve_Mins            float64
Night_Mins          float64
Intl_Mins           float64
CustServ_Calls      int64
Churn               object
Intl_Plan           object
Vmail_Plan          object
Day_Calls           int64
Day_Charge          float64
Eve_Calls           int64
Eve_Charge          float64
Night_Calls         int64
Night_Charge        float64
Intl_Calls          int64
Intl_Charge         float64
State              object
Area_Code           int64
Phone              object
dtype: object
```

```
telco["Intl_Plan"].value_counts()
```

```
no      3010
yes      323
Name: Intl_Plan, dtype: int64
```

```
telco["Intl_Plan"].replace({"yes":1, "no":0},inplace=True)
telco["Churn"].replace({"yes":1, "no":0},inplace=True)
telco["Vmail_Plan"].replace({"yes":1, "no":0},inplace=True)
```


Example Churn – Categorical to Dummies

```
df = pd.get_dummies(telco["State"])  
telco.drop("State",axis=1,inplace=True)  
df
```


	AK	AL	AR	AZ	CA	CO	CT	DC	DE	FL	...	SD	TN	TX	UT	VA	VT	WA	WI	WV	WY
0	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
2	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
3	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
4	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
...
3328	0	0	0	1	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
3329	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	1	0
3330	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
3331	0	0	0	0	0	0	1	0	0	0	...	0	0	0	0	0	0	0	0	0	0
3332	0	0	0	0	0	0	0	0	0	0	...	0	1	0	0	0	0	0	0	0	0

3333 rows × 51 columns

```
telco = pd.merge(telco,df,left_index=True,right_index=True,how="inner")
```

Example Churn – Correlations

```
plt.figure(figsize=(12, 8))
sns.heatmap(cbar=False, annot=True, fmt=".2f", data=telco.drop("State", axis=1).corr(), cmap='coolwarm')
plt.title('Correlation Matrix')
plt.show()
```



```
telco["Vmail_Message"].groupby(telco["Vmail_Plan"]).describe()
```

	count	mean	std	min	25%	50%	75%	max
Vmail_Plan								
0	2411.0	0.000000	0.000000	0.0	0.0	0.0	0.0	0.0
1	922.0	29.277657	7.559027	4.0	24.0	29.0	34.0	51.0

```
telco["Day_Mins"] / telco["Day_Charge"]
```

```
0      5.881961
1      5.882781
2      5.882069
3      5.882122
4      5.882145
...
3328    5.883239
3329    5.881904
3330    5.881588
3331    5.881706
3332    5.882058
Length: 3333, dtype: float64
```

```
telco["Total_Charge"] = telco["Day_Charge"] + telco["Eve_Charge"] + telco["Night_Charge"] + telco["Intl_Charge"]  
telco["Total_Mins"] = telco["Day_Mins"] + telco["Eve_Mins"] + telco["Night_Mins"] + telco["Intl_Mins"]  
telco["Avg_Rate"] = telco["Total_Charge"] / telco["Total_Mins"]  
telco["Avg_Rate"].describe()
```

```
count    3333.000000  
mean      0.100354  
std       0.008440  
min       0.066950  
25%       0.094893  
50%       0.100385  
75%       0.106056  
max       0.129791  
Name: Avg_Rate, dtype: float64
```

Example Churn – Dataset after preprocessing

```
telco.drop("Phone",axis=1,inplace=True)
```

```
print(telco.columns)
```

```
Index(['Account_Length', 'Vmail_Message', 'Day_Mins', 'Eve_Mins', 'Night_Mins',  
      'Intl_Mins', 'CustServ_Calls', 'Churn', 'Intl_Plan', 'Vmail_Plan',  
      'Day_Calls', 'Day_Charge', 'Eve_Calls', 'Eve_Charge', 'Night_Calls',  
      'Night_Charge', 'Intl_Calls', 'Intl_Charge', 'Area_Code',  
      'Total_Charge', 'Total_Mins', 'Avg_Rate', 'AK', 'AL', 'AR', 'AZ', 'CA',  
      'CO', 'CT', 'DC', 'DE', 'FL', 'GA', 'HI', 'IA', 'ID', 'IL', 'IN', 'KS',  
      'KY', 'LA', 'MA', 'MD', 'ME', 'MI', 'MN', 'MO', 'MS', 'MT', 'NC', 'ND',  
      'NE', 'NH', 'NJ', 'NM', 'NV', 'NY', 'OH', 'OK', 'OR', 'PA', 'RI', 'SC',  
      'SD', 'TN', 'TX', 'UT', 'VA', 'VT', 'WA', 'WI', 'WV', 'WY'],  
      dtype='object')
```

Example Churn – First Quick Classification

```
X = telco.drop("Churn",axis=1)
y = np.array(telco["Churn"])
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, stratify=y, random_state=30)
```

```
from sklearn.tree import DecisionTreeClassifier
tree = DecisionTreeClassifier(random_state=40)
tree.fit(X_train, y_train)
print("Accuracy on training set: {:.3f}".format(tree.score(X_train, y_train)))
print("Accuracy on test set: {:.3f}".format(tree.score(X_test, y_test)))
y_pred = tree.predict(X_test)
from sklearn.metrics import confusion_matrix, precision_score, recall_score
print("Precision on test set: {:.3f}".format(precision_score(y_test, y_pred)))
print("Recall on test set: {:.3f}".format(recall_score(y_test, y_pred)))
confusion_matrix(y_test, y_pred)
```

Accuracy on training set: 1.000


Accuracy on test set: 0.950

Precision on test set: 0.793

Recall on test set: 0.884

```
array([[685, 28],
       [ 14, 107]], dtype=int64)
```

Example Churn – First Quick Classification



Example Churn – Tuning Logistic Regression (1/7)

```
used_features = ['Account_Length', 'Vmail_Message', 'Day_Mins', 'Eve_Mins', 'Night_Mins',  
                'Intl_Mins', 'CustServ_Calls', 'Intl_Plan', 'Vmail_Plan', 'Day_Calls', 'Day_Charge',  
                'Eve_Calls', 'Eve_Charge', 'Night_Calls', 'Night_Charge', 'Intl_Calls', 'Intl_Charge',  
                'Area_Code', 'AK', 'AL', 'AR', 'AZ', 'CA', 'CO', 'CT', 'DC', 'DE', 'FL', 'GA', 'HI',  
                'IA', 'ID', 'IL', 'IN', 'KS', 'KY', 'LA', 'MA', 'MD', 'ME', 'MI', 'MN', 'MO', 'MS',  
                'MT', 'NC', 'ND', 'NE', 'NH', 'NJ', 'NM', 'NV', 'NY', 'OH', 'OK', 'OR', 'PA', 'RI',  
                'SC', 'SD', 'TN', 'TX', 'UT', 'VA', 'VT', 'WA', 'WI', 'WV', 'WY']
```

```
l_copies = [0]  
l_threshold = [0.5]  
l_penalty = ["l2"]  
l_C = [1,10,0.1]  
l_tol = [10**(-4)]  
l_max_iter = [10000]  
  
from sklearn.linear_model import LogisticRegression
```

Data not scaled

Best F1 and best parameters:

0.32499999999999996 (0, 0.5, 'l2', 1, 0.0001, 10000)

	Training	Validation	Test
Accuracy	0.87	0.865	0.892
Precision	0.639	0.591	0.75
Recall	0.245	0.224	0.375
F1	0.354	0.325	0.5
Confusion matrix Row 1	[1837, 44]	[666, 18]	[279, 6]
Confusion matrix Row 2	[241, 78]	[90, 26]	[30, 18]

Example Churn – Tuning Logistic Regression (2/7)

```
used_features = ['Account_Length', 'Vmail_Message', 'Day_Mins', 'Eve_Mins', 'Night_Mins',  
                'Intl_Mins', 'CustServ_Calls', 'Intl_Plan', 'Vmail_Plan', 'Day_Calls', 'Day_Charge',  
                'Eve_Calls', 'Eve_Charge', 'Night_Calls', 'Night_Charge', 'Intl_Calls', 'Intl_Charge',  
                'Area_Code', 'AK', 'AL', 'AR', 'AZ', 'CA', 'CO', 'CT', 'DC', 'DE', 'FL', 'GA', 'HI',  
                'IA', 'ID', 'IL', 'IN', 'KS', 'KY', 'LA', 'MA', 'MD', 'ME', 'MI', 'MN', 'MO', 'MS',  
                'MT', 'NC', 'ND', 'NE', 'NH', 'NJ', 'NM', 'NV', 'NY', 'OH', 'OK', 'OR', 'PA', 'RI',  
                'SC', 'SD', 'TN', 'TX', 'UT', 'VA', 'VT', 'WA', 'WI', 'WV', 'WY']
```

```
l_copies = [0]  
l_threshold = [0.5]  
l_penalty = ["l2"]  
l_C = [1, 10, 0.1]  
l_tol = [10**(-4)]  
l_max_iter = [10000]  
  
from sklearn.linear_model import LogisticRegression
```

Data scaled

Best F1 and best parameters:

0.3312101910828026 (0, 0.5, 'l2', 1, 0.0001, 10000)

	Training	Validation	Test
Accuracy	0.862	0.869	0.874
Precision	0.577	0.634	0.688
Recall	0.188	0.224	0.229
F1	0.284	0.331	0.344
Confusion matrix Row 1	[1837, 44]	[669, 15]	[280, 5]
Confusion matrix Row 2	[259, 60]	[90, 26]	[37, 11]

Example Churn – Tuning Logistic Regression (3/7)

```
used_features = ['Account_Length', 'Vmail_Message', 'Day_Mins', 'Eve_Mins', 'Night_Mins',  
                'Intl_Mins', 'CustServ_Calls', 'Intl_Plan', 'Vmail_Plan', 'Day_Calls', 'Day_Charge',  
                'Eve_Calls', 'Eve_Charge', 'Night_Calls', 'Night_Charge', 'Intl_Calls', 'Intl_Charge']
```

```
l_copies = [0]  
l_threshold = [0.5]  
l_penalty = ["l2"]  
l_C = [1,10,0.1]  
l_tol = [10**(-4)]  
l_max_iter = [10000]  
  
from sklearn.linear_model import LogisticRegression
```

Data scaled

Best F1 and best parameters:

0.3312101910828026 (0, 0.5, 'l2', 1, 0.0001, 10000)

	Training	Validation	Test
Accuracy	0.862	0.869	0.874
Precision	0.577	0.634	0.688
Recall	0.188	0.224	0.229
F1	0.284	0.331	0.344
Confusion matrix Row 1	[1837, 44]	[669, 15]	[280, 5]
Confusion matrix Row 2	[259, 60]	[90, 26]	[37, 11]

Example Churn – Tuning Logistic Regression (4/7)

```
used_features = ['Account_Length', 'Vmail_Message', 'Day_Mins', 'Eve_Mins', 'Night_Mins',  
                'Intl_Mins', 'CustServ_Calls', 'Intl_Plan', 'Day_Calls', 'Eve_Calls', 'Night_Calls',  
                'Intl_Calls']
```

```
l_copies = [0]  
l_threshold = [0.5]  
l_penalty = ["l2"]  
l_C = [1, 10, 0.1]  
l_tol = [10**(-4)]  
l_max_iter = [10000]  
  
from sklearn.linear_model import LogisticRegression
```

Data scaled

Best F1 and best parameters:

0.3312101910828026 (0, 0.5, 'l2', 1, 0.0001, 10000)

	Training	Validation	Test
Accuracy	0.862	0.869	0.874
Precision	0.577	0.634	0.688
Recall	0.188	0.224	0.229
F1	0.284	0.331	0.344
Confusion matrix Row 1	[1837, 44]	[669, 15]	[280, 5]
Confusion matrix Row 2	[259, 60]	[90, 26]	[37, 11]

Example Churn – Tuning Logistic Regression (5/7)

```
used_features = ['Account_Length', 'Vmail_Message', 'Day_Mins', 'Eve_Mins', 'Night_Mins',  
                'Intl_Mins', 'CustServ_Calls', 'Intl_Plan', 'Day_Calls', 'Eve_Calls', 'Night_Calls',  
                'Intl_Calls', 'Total_Charge', 'Avg_Rate']
```

```
l_copies = [0]  
l_threshold = [0.5]  
l_penalty = ["l2"]  
l_C = [100, 10, 1000, 1]  
l_tol = [10**(-4)]  
l_max_iter = [10000]  
  
from sklearn.linear_model import LogisticRegression
```

Data scaled

Best F1 and best parameters:

0.4311377245508982 (0, 0.5, 'l2', 100, 0.0001, 10000)

	Training	Validation	Test
Accuracy	0.871	0.881	0.898
Precision	0.628	0.706	0.75
Recall	0.27	0.31	0.438
F1	0.377	0.431	0.553
Confusion matrix Row 1	[1830, 51]	[669, 15]	[278, 7]
Confusion matrix Row 2	[233, 86]	[80, 36]	[27, 21]

```
used_features = ['Account_Length', 'Vmail_Message', 'Day_Mins', 'Eve_Mins', 'Night_Mins',  
                'Intl_Mins', 'CustServ_Calls', 'Intl_Plan', 'Day_Calls', 'Eve_Calls', 'Night_Calls',  
                'Intl_Calls', 'Total_Charge', 'Avg_Rate']
```

```
l_copies = [0]  
l_threshold = [0.5, 0.4, 0.3, 0.2, 0.1]  
l_penalty = ["l2"]  
l_C = [100, 10, 1000, 1]  
l_tol = [10**(-4)]  
l_max_iter = [10000]  
  
from sklearn.linear_model import LogisticRegression
```

Data scaled

Best F1 and best parameters:

0.5409252669039145 (0, 0.2, 'l2', 10, 0.0001, 10000)

	Training	Validation	Test
Accuracy	0.827	0.839	0.85
Precision	0.438	0.461	0.486
Recall	0.68	0.655	0.75
F1	0.533	0.541	0.59
Confusion matrix Row 1	[1603, 278]	[595, 89]	[247, 38]
Confusion matrix Row 2	[102, 217]	[40, 76]	[12, 36]

Example Churn – Tuning Logistic Regression (7/7)

```
used_features = ['Account_Length', 'Vmail_Message', 'Day_Mins', 'Eve_Mins', 'Night_Mins',  
                'Intl_Mins', 'CustServ_Calls', 'Intl_Plan', 'Day_Calls', 'Eve_Calls', 'Night_Calls',  
                'Intl_Calls', 'Total_Charge', 'Avg_Rate']
```

```
l_copies = [0,1,2,3,4]  
l_threshold = [0.5,0.4,0.3,0.2,0.1]  
l_penalty = ["l2"]  
l_C = [100,10,1000,1]  
l_tol = [10**(-4)]  
l_max_iter = [10000]  
  
from sklearn.linear_model import LogisticRegression
```

Data scaled

Best F1 and best parameters:
0.5625 (1, 0.4, 'l2', 10, 0.0001, 10000)

	Training	Validation	Test
Accuracy	0.815	0.86	0.862
Precision	0.646	0.514	0.515
Recall	0.592	0.621	0.708
F1	0.618	0.562	0.596
Confusion matrix Row 1	[1674, 207]	[616, 68]	[253, 32]
Confusion matrix Row 2	[260, 378]	[44, 72]	[14, 34]

- (1) Go through the notebook and try to roughly understand the code behind the previous slides
- (2) Adapt the last cells of the workbook and tune the parameters for
 - a) Kernelized Support Vectors Machines
 - b) Random Forest
 - c) Gradient Boosted Trees
- (3) If you still have time: Create new cells in the notebook for another model (e.g. neural network, k-nearest neighbors, decision tree) and tune appropriate model parameters
- (4) Which model and which parameter setting results in the best F1-score on the validation set? What is the corresponding F1-score on the test set? Which model / setting would you recommend to use for predictions on unknown data?

Discuss the course. What went well, what can be approved?

Align on the top 2 stars and wishes within your group, i.e. present

- 2 things that worked very well (due to the course design, but could also be best practices that you applied as team or individuals)
- 2 concrete suggestions how the course design can be improved (please be very concrete and specific)

Focus on the 2 most important ones in each category. Sort out and rank all of your ideas